






(a) Three-branch feature extraction



(b) CrossGate module



(c) Classification head

B denotes batch size. The two MaxViT encoders have independent weights. Labels: 0 = non-glaucoma, 1 = glaucoma. Dashed arrows: post-hoc test calibration; T is fitted on validation logits. Training uses class-weighted cross-entropy on raw logits.