






(a) ResNeXt3D encoder



(b) Executed residual block



(c) MaxViT-Tiny encoder (each view)

Each MaxViT block: MBConv, window attention + FFN, grid attention + FFN. The two encoders have independent weights.
Executed graph: no GroupNorm after the final 1x1x1 convolution in the residual branch (g3 is declared but unused).